Homework 2

Theory [0.8 pt each]

Unless otherwise stated,

- all 1-D vectors, such as x and w below, are column vectors by default.
- the classification problem is binary.
- a lower case letter variable is a scalar or vector, where a upper case letter (in any font) is a matrix.

1. Given a sample with feature vector $x = [1.1, 2.2, 3.3]^T$, what is its augmented feature vector?

2. If the weight vector of a linear classifier is $w = [1, 0, 1, 0]^T$, and we define that a sample belongs to class $+1$ if $w^T x > 0$ and -1 if $w^T x < 0$ where x is the augmented feature vector of the sample, what is the class of the sample?

3. When discussing the sum of error squares loss function in the class, we used augmented but not normalized augmented (normalized and augmented) feature vectors. Please rewrite that loss function in terms of **normalized augmented** feature vectors. Let x_i'' be the normalized augmented feature vector of the i-th sample, and w be the weight vector of the classifier. A correct prediction shall satisfy $w^T x_i'' > 0$ regardless of the class of the sample because x_i'' has been normalized. You may use a computational algebra system to help – but it is not required. It might be easier by hand on scratch paper.

4. Please find the solution for minimizing the new loss function. Keep variables and font style consistent with those in the class notes/slides. Denote $\mathbb{X}$ as a matrix, each **row (not column as in slides)** of which corresponds to a sample and each column of which corresponds to a feature value except the rightmost one (i.e., the bias term).

5. Everyone gets free point for this problem.

Programming

6. [3pts] The demo `2_Linear_classifiers.ipynb` provides snippets of code to achieve a linear classifier using the minimization of the sum of the squared errors, and the visualization of the classifier. Now please define a function `plot_mse` that visualizes the training samples and the classifier on the sample plot.

   Inputs/arguments:

   - `X`: a 2-D numpy array such that `X[i]`, which is 1-D numpy array, corresponds to the i-th sample's feature vector (not augmented). The shape of `X[i]` is 1-by-2.
   - `y`: a 1-D numpy array such that `y[i]`, which is a scalar $\in [+1, -1]$, is the label of the i-th sample.
   - `filename`: a string, the path to save the plot.

   Output/return:

   - `w`: the weight vector of the classifier, where the 1st element corresponds to the 1st dimension of a feature vector, the 2nd element corresponds to the 2nd dimension of a feature vector, and the last term is the bias.

7. [3pts] Redo the function above using Fisher's linear discriminant. Save the function as `plot_fisher`.

   Note that on the slides, when discussing Fisher's, the **w** does not contain the bias. But when returning **w** here, please add the bias as the last element.

Specifications regarding the plots

In your 2-D plots, let the horizontal axis correspond to the first/leftmost feature column in `X`, and the vertical axis correspond to the 2nd feature column in `X`.

For plot styles and properties, please plot class $+1$ samples in blue dots and class -1 samples in red dots. For all other objects and properties (e.g., line width, line color, marker size, marker border color, etc.), use default settings.

To compute and plot the decision hyperplane, compute two points and draw a line connecting them. Let the horizontal coordinates of the two points be

```
x_ticks = numpy.array([numpy.min(X[:,0]), numpy.max(X[:,0])])
```

then their vertical coordinates are

```
y_ticks = (x_ticks * w[0] +x[2])/w[1]
```

Finally, just plot

```
matplotlib.pyplot.plot(x_ticks, y_ticks)
```

Be sure that the plot range does not go beyond the sample range. To do that, use `xlim` and `ylim` of matplotlib to set:

```
matplotlib.pyplot.xlim(numpy.min(X[:,0]), numpy.max(X[:,0]))
matplotlib.pyplot.ylim(numpy.min(X[:,1]), numpy.max(X[:,1]))
```

An example plot is given below (may float to next page in pdf):
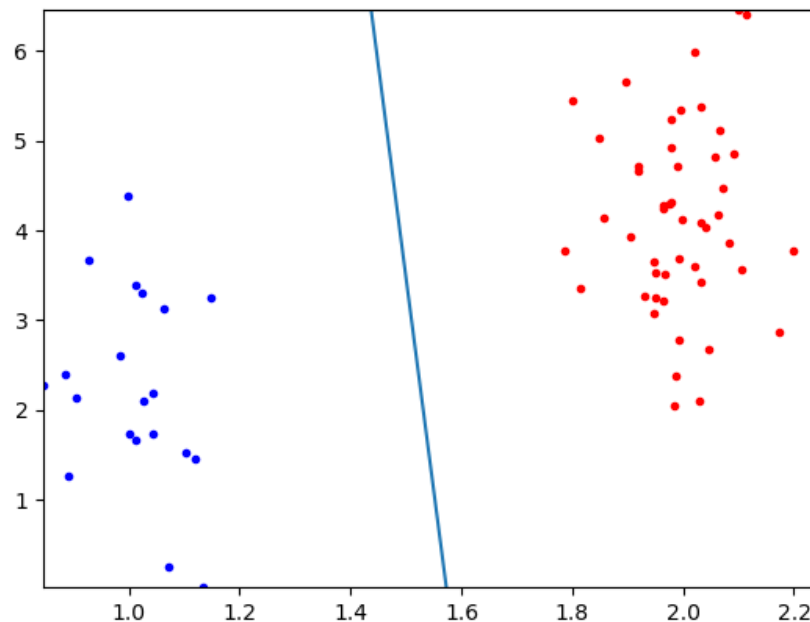

Figure 1: example plot

How to submit

For theory part, submit as a PDF file. For programming part, work on the template `hw2.py` and submit it.

How to view this in nice PDF

```
pandoc -s hw2.md -o hw2.pdf
```

Programming hints

- Note in programming assignment, each sample is a row in `X` instead of a column vector as in the math in slides.

- We will use the following example X and y to explain.

```
X = numpy.array([[1,2,3], # class 1
                 [4,5,6], # class -1
                 [7,8,9]]) # class 1
y = numpy.array([1,-1,1])
```

The `X` is NOT augmented nor normalized.

- How to separate samples of the two classes?

```
X1 = X[y == +1] # samples of class 1
X2 = X[y == -1] # samples of class 2
```

- How to compute $\mathbf{m}_i$?

Use Class 1 as example:

```
m1 = numpy.mean(X1, axis=0)
```

If using `axis=1`, then it computes mean per row instead of mean per column.

- How to know the size of a class $|C_i|$?

There are many ways, you could measure the number of rows for samples in the class like `X1.shape[0]`, or you could measure the number of elements in `y` equal to i like `numpy.sum(y==1)`.

- How to plot with proper range? Search `x_lim` in this doc to see the code. Just use that code.

- How to visualize the hyperplane from $\mathbf{w}$? Search `x_ticks` in this doc to see the code. Just use that code.

Additional programming hints for Fisher's

- How to compute $\mathbf{S}_i$? Let's first see how to transform the sum of dot products into matrix operation. $\mathbf{m}_1 = [4, 5, 6]^T$. Using class 1 as the example, by the definition,

$$\mathbf{S}_1 = \left[\begin{pmatrix} 1 \\ 2 \\ 3 \end{pmatrix} - \begin{pmatrix} 4 \\ 5 \\ 6 \end{pmatrix} \right] \left[\begin{pmatrix} 1 \\ 2 \\ 3 \end{pmatrix} - \begin{pmatrix} 4 \\ 5 \\ 6 \end{pmatrix} \right]^T + \left[\begin{pmatrix} 7 \\ 8 \\ 9 \end{pmatrix} - \begin{pmatrix} 4 \\ 5 \\ 6 \end{pmatrix} \right] \left[\begin{pmatrix} 7 \\ 8 \\ 9 \end{pmatrix} - \begin{pmatrix} 4 \\ 5 \\ 6 \end{pmatrix} \right]^T \tag{1}$$

$$= \left[\begin{pmatrix} 1 \\ 2 \\ 3 \end{pmatrix} - \begin{pmatrix} 4 \\ 5 \\ 6 \end{pmatrix}, \begin{pmatrix} 7 \\ 8 \\ 9 \end{pmatrix} - \begin{pmatrix} 4 \\ 5 \\ 6 \end{pmatrix} \right] \left[\begin{pmatrix} 1 \\ 2 \\ 3 \end{pmatrix} - \begin{pmatrix} 4 \\ 5 \\ 6 \end{pmatrix}, \begin{pmatrix} 7 \\ 8 \\ 9 \end{pmatrix} - \begin{pmatrix} 4 \\ 5 \\ 6 \end{pmatrix} \right]^T \tag{2}$$

$$= \left[\begin{pmatrix} 1 & 7 \\ 2 & 8 \\ 3 & 9 \end{pmatrix} - \begin{pmatrix} 4 & 4 \\ 5 & 5 \\ 6 & 6 \end{pmatrix} \right] \left[\begin{pmatrix} 1 & 7 \\ 2 & 8 \\ 3 & 9 \end{pmatrix} - \begin{pmatrix} 4 & 4 \\ 5 & 5 \\ 6 & 6 \end{pmatrix} \right]^T \tag{3}$$

$$= \begin{pmatrix} 1-4 & 7-4 \\ 2-5 & 8-5 \\ 3-6 & 9-6 \end{pmatrix} \begin{pmatrix} 1-4 & 7-4 \\ 2-5 & 8-5 \\ 3-6 & 9-6 \end{pmatrix}^T \tag{4}$$

$$= \begin{pmatrix} 18 & 18 & 18 \\ 18 & 18 & 18 \\ 18 & 18 & 18 \end{pmatrix} \tag{5}$$

From this example (particularly Eq.~(3)) we can see that,

$$\mathbf{S}_i = \sum_{\mathbf{x} \in C_i} (\mathbf{x} - \mathbf{m}_i)(\mathbf{x} - \mathbf{m}_i)^T$$

can be rewritten into the multiplication between two matrixes:

$$\underbrace{\mathbf{S}_i}_{d \times d} = \underbrace{(\mathbf{X_i} - \mathbf{M_i})}_{d \times |C_i|} (\mathbf{X_i} - \mathbf{M_i})^T$$

where $\mathbf{X}_i = \begin{pmatrix} | & & | & & | \\ \mathbf{x}_{i_1} & \cdots & \mathbf{x}_{i_j} & \cdots & \mathbf{x}_{i_{|C_i|}} \\ | & & | & & | \end{pmatrix}$ comprises of vertically stacked $|C_i|$ samples $[\mathbf{x}_{i_1}, \mathbf{x}_{i_2}, \cdots, \mathbf{x}_{i_j}, \mathbf{x}_{i_{|C_i|}}]$

in class i (a sample per column) and $\mathbf{M}_i = \underbrace{\begin{pmatrix} | & & | & | & | \\ \mathbf{m}_i & \cdots & \mathbf{m}_i & \cdots & \mathbf{m}_i \\ | & & | & | & | \end{pmatrix}}_{|C_i| \text{ columns}}$ duplicates the mean vector $\mathbf{m}_i$

by $|C_i|$ times vertically, and d is the number of features.

- How to programmingly construct the "duplication" matrix $\mathbf{M_i}$?

The example below duplicates a 1-D vector row-wisely. It's not the exact solution but shall give you the idea.

```
>>> m1=numpy.array([4,5,6])
>>> numpy.array([m1]*2)
array([[4, 5, 6],
       [4, 5, 6]])
```